

leakage

April 28, 2025

```
[10]: # 0. Imports & Logging
import logging
from pathlib import Path
import pandas as pd
import spacy
from rapidfuzz import fuzz
from collections import Counter
from IPython.display import display
import matplotlib.pyplot as plt
from config import OUTPUT_FOLDER
```

```
[11]: # %%
# 9. Summary by noun_gender × adjective_gender (with percentages)
# 1. Configuration
CSV_IN = Path(OUTPUT_FOLDER) / "sentences_final.csv"

# %%
# 2. Preflight & Load
if not CSV_IN.exists():
    logger.error(f"Input file not found: {CSV_IN}")
    raise FileNotFoundError(f"{CSV_IN} does not exist")
logger.info(f"Reading input CSV from {CSV_IN}")
df = pd.read_csv(CSV_IN, encoding=ENCODING, on_bad_lines="warn")
```

```
2025-04-28 10:22:16 INFO      sentence_cleaner: Reading input CSV from
/Users/matthijstentije/University/MSc_Data-
Science/Thesis/MSc_Data_Science_Thesis/phase_02/output/sentences_final.csv
```

```
[12]: # -----
# 7. Leakage summary
# -----
pct_multi = (df['total_in_lists'] > 1).mean() * 100
# -----
# 8. Print example sentences with both male & female adjectives
# -----
# Print to console (and log)
print(f"% sentences with more than one adjective: {pct_multi:.2f}%")
```

```
logger.info("Leakage summary - >1 count: %.2f%%;",
            pct_multi)
```

```
2025-04-28 10:22:17 INFO      sentence_cleaner: Leakage summary - >1 count:
36.36%;
```

```
% sentences with more than one adjective: 36.36%
```

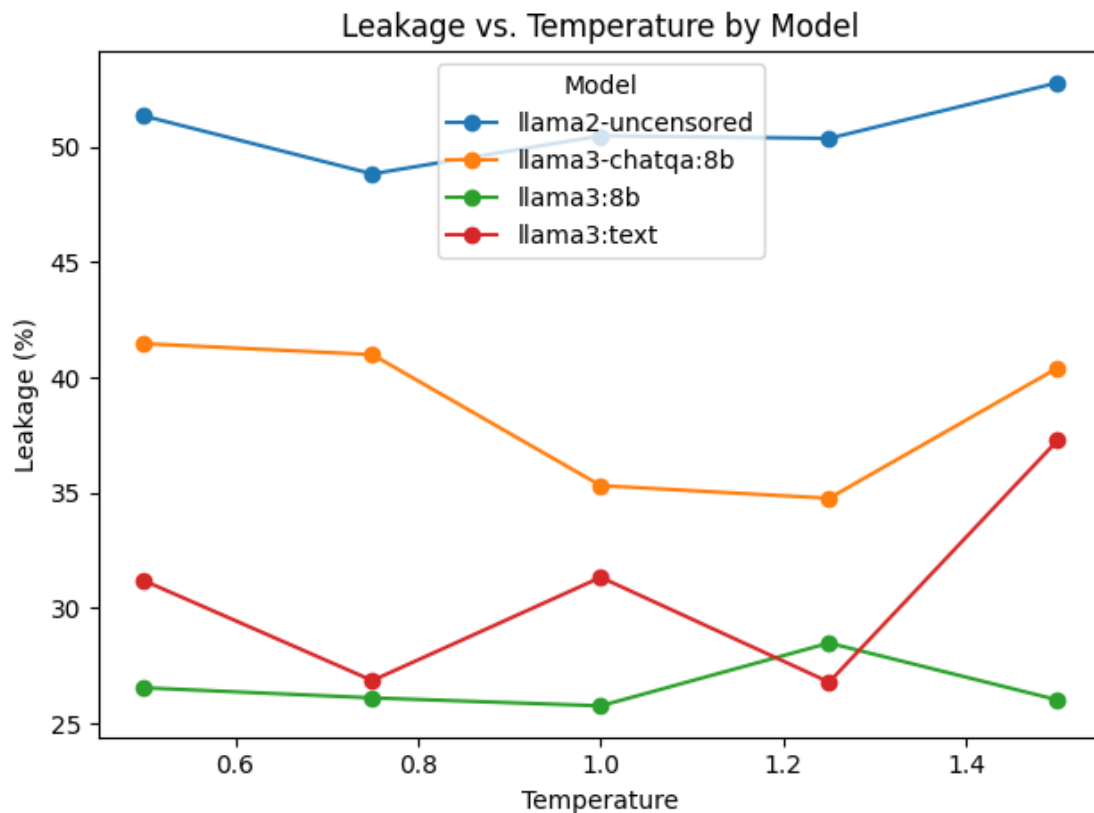
```
[13]: # Aggregate leakage by model × temperature
agg_model_temp = (
    df
    .groupby(['model', 'temperature'])
    .apply(lambda sub: (sub['total_in_lists'] > 1).mean() * 100)
    .reset_index(name='pct_multi')
)

# Plot each model in its own color
plt.figure()
for model in agg_model_temp['model'].unique():
    subset = agg_model_temp[agg_model_temp['model'] == model]
    plt.plot(
        subset['temperature'],
        subset['pct_multi'],
        marker='o',
        label=model
    )

plt.xlabel('Temperature')
plt.ylabel('Leakage (%)')
plt.title('Leakage vs. Temperature by Model')
plt.legend(title="Model")
plt.tight_layout()
plt.show()
```

```
/var/folders/v_/ypb5b1nx079dj_hcjyz0t47c0000gn/T/ipykernel_61298/3840618518.py:5
: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.
```

```
.apply(lambda sub: (sub['total_in_lists'] > 1).mean() * 100)
```



```
[17]: print(df)
```

	word	sentence \
0	lamme	de lamme man was een vrolijke kerel
1	impopulair	de impopulaire jongen heeft het niet gemakkelijk
2	overmoedig	overmoedig is de man die je bent
3	vooraanstaand	hij was vooraanstaand in het leven van een man...
4	gewetenloos	deze gewetenloze man mannen jongen jongens hee...
...
15162	hitsig	het hitsige uitstralen van deze meisjes maakt ...
15163	glamoureuus	de glamour uitstraling van deze dame zorgt erv...
15164	bloedmooie	bloedmooie meisjes hebben een eigen unieke sfe...
15165	exotisch	exotische dames tonen altijd een speciale aanb...
15166	spichtig	een spichtig meisje is altijd mooi te vinden m...

	model	noun_gender	adjective_gender	temperature \
0	llama3:text	male	male	0.5
1	llama3:text	male	male	0.5
2	llama3:text	male	male	0.5
3	llama3:text	male	male	0.5
4	llama3:text	male	male	0.5

...
15162	llama2-uncensored	female	female	1.0
15163	llama2-uncensored	female	female	1.0
15164	llama2-uncensored	female	female	1.0
15165	llama2-uncensored	female	female	1.0
15166	llama2-uncensored	female	female	1.0

	total_runs	male_count	female_count	\
0	75	1	1	
1	75	1	0	
2	75	1	0	
3	75	1	0	
4	75	8	0	
...	
15162	55	1	1	
15163	55	0	1	
15164	55	0	1	
15165	55	0	1	
15166	55	0	1	

	male_matches	female_matches	\
0	['lamme']	['voorlijk']	
1	['impopulair']	[]	
2	['overmoedig']	[]	
3	['vooraanstaand']	[]	
4	['goddeloos', 'planmatig', 'incompetent', 'gew...	[]	
...	
15162	['impopulair']	['hitsig']	
15163	[]	['glamoureuus']	
15164	[]	['bloedmooie']	
15165	[]	['exotisch']	
15166	[]	['spichtig']	

	total_in_lists	prompt_type
0	2	M→M
1	1	M→M
2	1	M→M
3	1	M→M
4	8	M→M
...
15162	2	F→F
15163	1	F→F
15164	1	F→F
15165	1	F→F
15166	1	F→F

[15167 rows x 13 columns]

```
[14]: # --- 1. derive prompt_type ---
# Map "male"→"M", "female"→"F" and build e.g. "M→F", "F→M", etc.
gender_map = {"male": "M", "female": "F"}
df["prompt_type"] = (
    df["adjective_gender"].map(gender_map)
    + "→"
    + df["noun_gender"].map(gender_map)
)

# --- 3. Prompt-structure interaction ---
pt = (
    df
    .groupby("prompt_type")
    .apply(lambda sub: pd.Series({
        "pct_multi": (sub["total_in_lists"]>1).mean()*100,
    }))
    .round(2)
)
print("\nLeakage by prompt structure:\n", pt)


import numpy as np
# grouped bar chart
labels = pt.index.tolist()
x = np.arange(len(labels))
plt.bar(x-0.15, pt["pct_multi"], width=0.3, label="Multi-adj")
plt.xticks(x, labels); plt.ylabel("Leakage (%)")
plt.title("Leakage by Prompt Structure")
plt.legend(); plt.tight_layout()
plt.show()
```

Leakage by prompt structure:

	pct_multi
prompt_type	
F→F	37.12
F→M	35.29
M→F	36.04
M→M	36.98

/var/folders/v/_ypb5b1nx079dj_hcjyz0t47c0000gn/T/ipykernel_61298/3785049157.py:14: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated, and in a future version of pandas the grouping columns will be excluded from the operation. Either pass `include_groups=False` to exclude the groupings or explicitly select the grouping columns after groupby to silence this warning.

```
.apply(lambda sub: pd.Series({
```



```
[15]: # leakage by prompt_type
pt = df.groupby('prompt_type').apply(lambda sub: pd.Series({
    'pct_multi': (sub['total_in_lists']>1).mean()*100,
    'pct_co':    ((sub['male_count']>0)&(sub['female_count']>0)).mean()*100
})).round(2)

print(pt)

import numpy as np

labels = pt.index.tolist()
multi = pt['pct_multi'].values
co     = pt['pct_co'].values
x = np.arange(len(labels))


plt.figure(figsize=(6,4))
plt.bar(x-0.15, multi, width=0.3, label='Multi-adj')
plt.bar(x+0.15, co,    width=0.3, label='Co-occur')
plt.xticks(x, labels)
plt.ylabel('Leakage (%)')
```

```
plt.title('Leakage by Prompt Structure')
plt.legend()
plt.tight_layout()
plt.show()
```

	pct_multi	pct_co
prompt_type		
F→F	37.12	11.62
F→M	35.29	10.77
M→F	36.04	10.62
M→M	36.98	10.21

```
/var/folders/v_/ypb5b1nx079dj_hcjyz0t47c0000gn/T/ipykernel_61298/3107383801.py:2
: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.
```

```
pt = df.groupby('prompt_type').apply(lambda sub: pd.Series({
```



```
[16]: # 1. Count each category
count_MM = len(df[(df['noun_gender']=='male') &
↪(df['adjective_gender']=='male')])
```

```

count_FF = len(df[(df['noun_gender']=='female') &
    ↪(df['adjective_gender']=='female')])
count_MF = len(df[(df['noun_gender']=='male') &
    ↪(df['adjective_gender']=='female')])
count_FM = len(df[(df['noun_gender']=='female') &
    ↪(df['adjective_gender']=='male')])

# 2. Totals
total_S = count_MM + count_FF
total_C = count_MF + count_FM
grand_total = len(df)

# 3. Percentages (relative to each half of the dataset)
pct_MM = count_MM / total_S * 100
pct_FF = count_FF / total_S * 100
pct_MF = count_MF / total_C * 100
pct_FM = count_FM / total_C * 100

# 4. Build summary table
summary = pd.DataFrame({
    '': [
        'Consistent with gender stereotype (S)',
        'Contradictory to gender stereotype (S)',
        'Total'
    ],
    '#MM': [f"{count_MM} ({pct_MM:.1f}%)", f"{count_MF} ({pct_MF:.1f}%)", ''],
    '#FF': [f"{count_FF} ({pct_FF:.1f}%)", f"{count_FM} ({pct_FM:.1f}%)", ''],
    '#Total': [f"{total_S} ({total_S/grand_total*100:.1f}%)",
        f"{total_C} ({total_C/grand_total*100:.1f}%)",
        f"{grand_total}"]
})

# 5. Display the table
print("Table 2: Labeling details with size & distribution", summary)

```

Table 2: Labeling details with size & distribution

#MM	#FF \		
0	Consistent with gender stereotype (S)	3789 (49.9%)	3804 (50.1%)
1	Contradictory to gender stereotype (S)	3789 (50.0%)	3785 (50.0%)
2	Total		
	#Total		
0	7593 (50.1%)		
1	7574 (49.9%)		
2	15167		